



Data Storage & Infrastructure

Amanda Farghli, Fatoumata
Drammeh, Khadiza Khanom,
Maisha Islam, Alexandr
Voronovich, Sharmin Zaman



Relational Database Systems

STRUCTURED

- Data stored in tables
- Not difficult to analyze



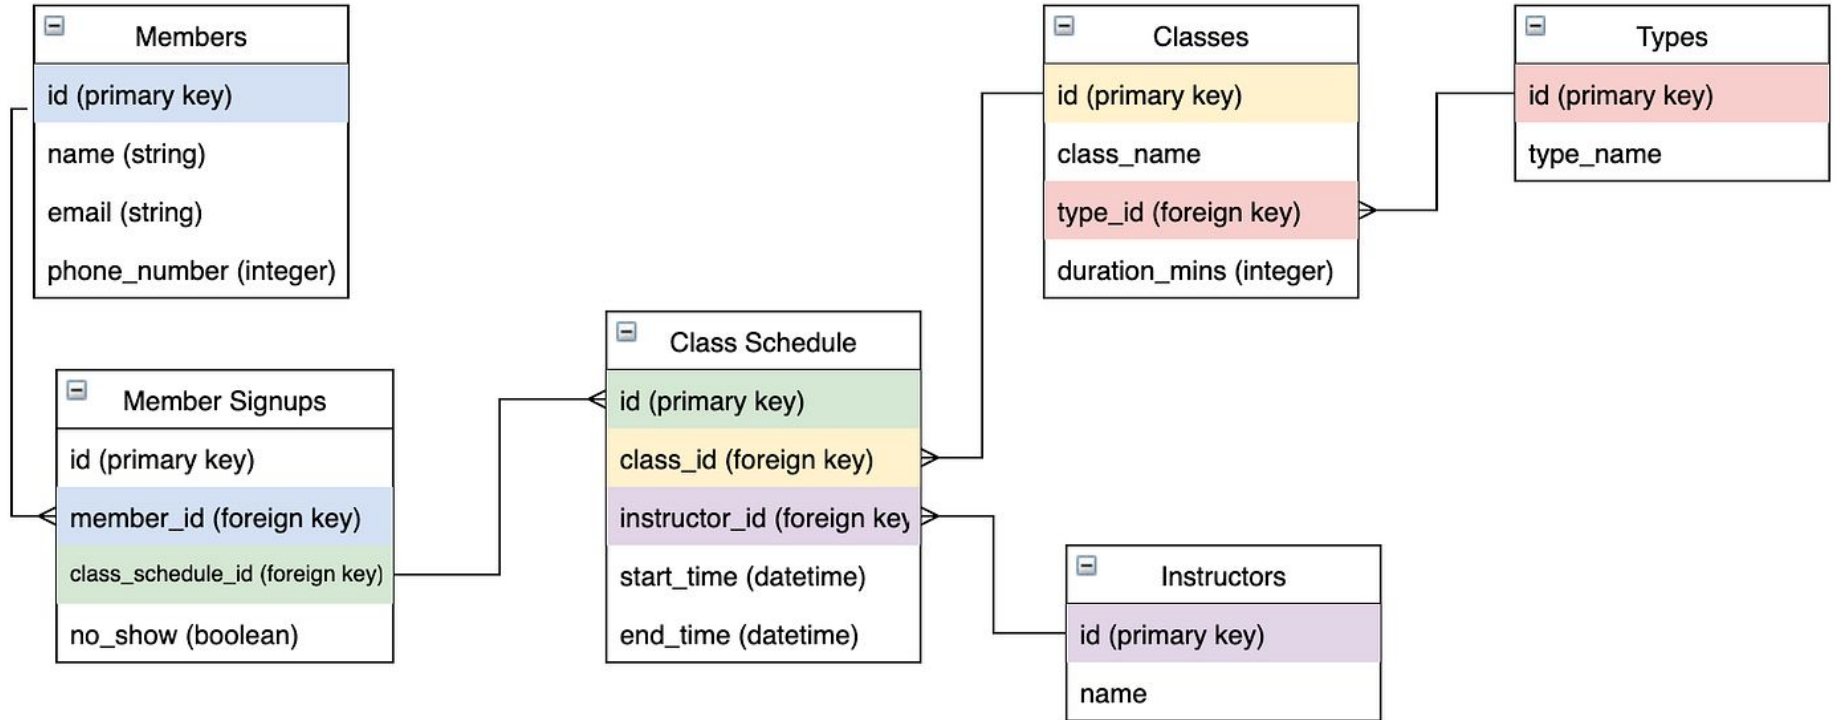
A relational database system follows a more rigid, structured model.

The database is fractured into various tables with an identifying key which defines data elements and their relationships to each other.


Key	Value
Name	John
ID_number	12345



Relational Database Systems



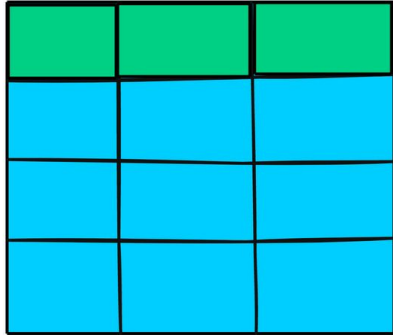
MYSQL & Postgre



PostgreSQL

SQL VS NoSQL


SQL




Relational

blog.algomaster.io

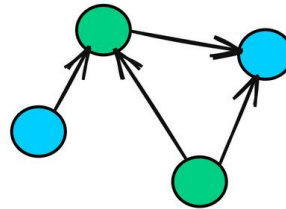
NoSQL



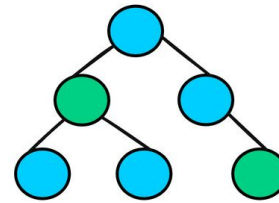
Key-Value



Column Store



Graph



Document


Document-oriented databases



Key Value

Key value database storage

Key	Value
customer:1:name	John Drake
customer:1:email	john.drake@gmail.com
customer:1:dob	24/11/1982
customer:1:mobile	7843241098

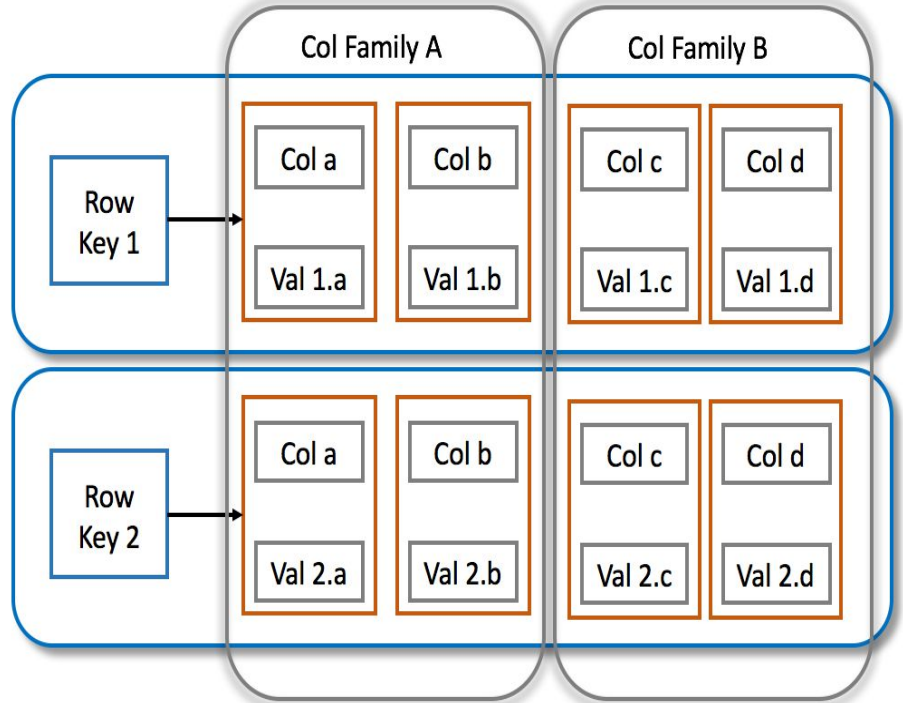


Amazon
DynamoDB

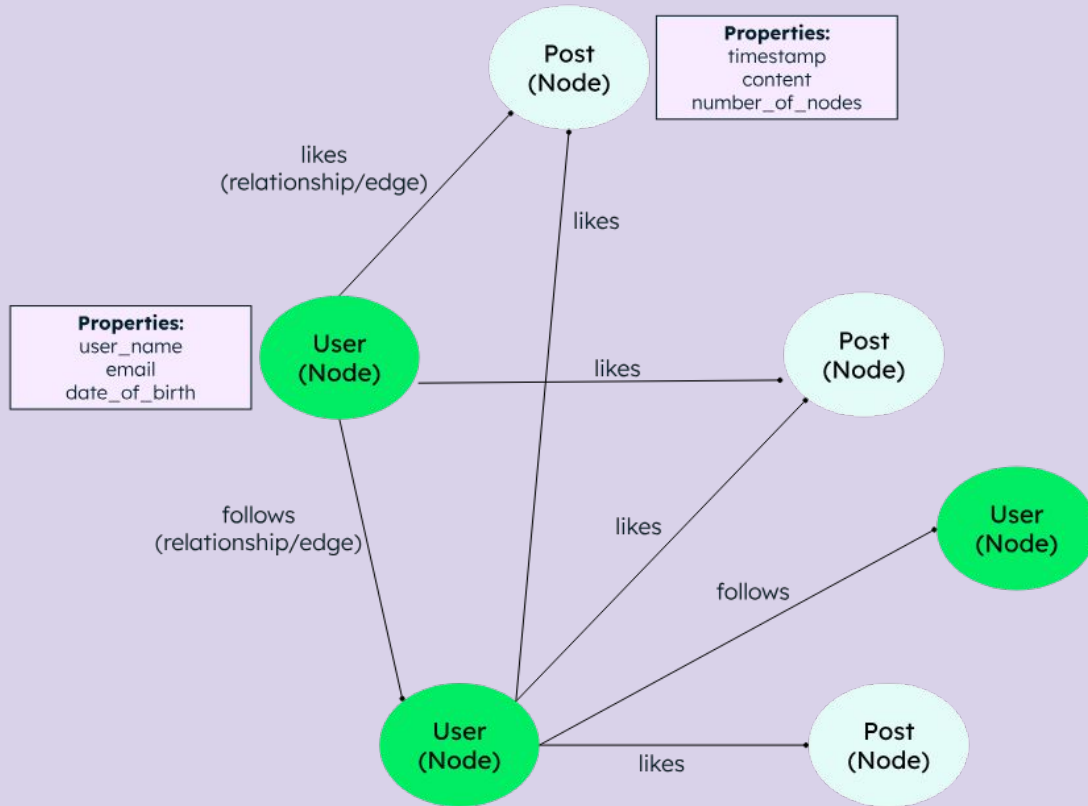


redis

Wide-column stores




Graph databases



NoSQL

Unstructured



Dynamic schema

Non relational model

Horizontally scalable

Suited for
hierarchical data storage

Which Database to Use & When

Relational Databases (SQL)

- Examples: MySQL, PostgreSQL, Oracle, SQL Server

NoSQL Databases

- Examples: MongoDB, CouchDB (JSON-like documents)

Relational Databases (SQL)

- Use for structured data with predictable schemas, like customer records or financial transactions. Ideal for complex queries, reporting, or apps needing strong consistency

NoSQL Databases


- Use for unstructured or semi-structured data, rapid scaling, or flexible schemas. Choose document stores (MongoDB) for content management, key-value (Redis) for caching, column-family (Cassandra) for big data

Elasticsearch & Cassandra



elasticsearch

- Full-text search & real-time analytics
- JSON-based, automatically indexed data
- Part of ELK stack (Logstash, Kibana)


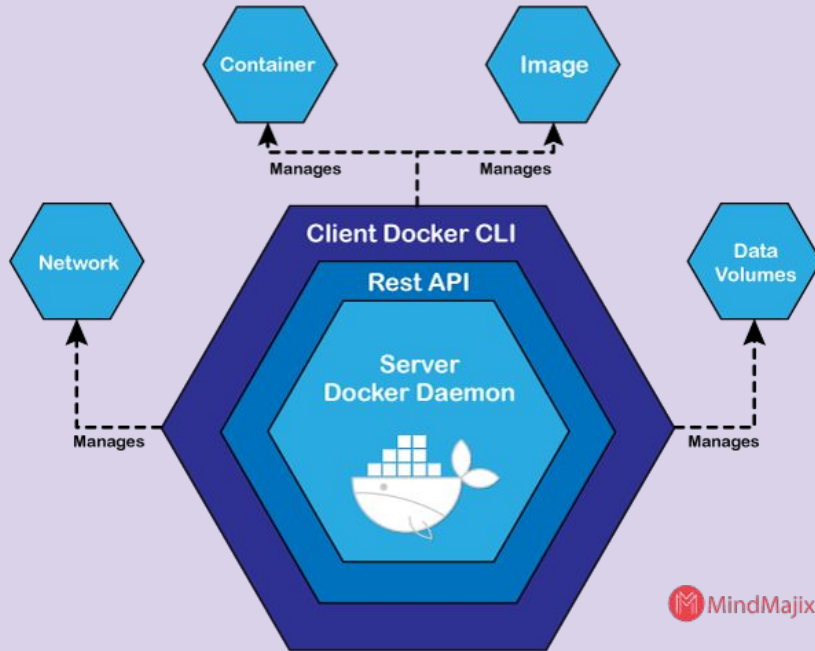


cassandra

- Distributed, column-based NoSQL database
- Built for scalability & uptime
- Handles massive data across servers

What is Docker?

A platform for developing, packaging, and running applications inside containers



docker

What are the uses for Docker?



**Collaboration
& Team
Development**



**Connecting
Databases**



Flexibility



Containers

How to Create a Container, push and pull an Image(MySQL)

- 1 Create a Container:** `docker run --name containerName -e MYSQL_ROOT_PASSWORD=password -p 3306:3306 -d mysql`
- 2 Make edits to container:** `CREATE DATABASE project; USE project; CREATE TABLE students (id INT...);`
- 3 Make a repository:** Got to <https://hub.docker.com> and create a repository and give it a name.
- 4 Build & Tague the Image:** `docker build -t yourUserName/repoName:tagName .`
- 5 Push to repository in Docker Hub:** `docker push yourUserName/repoName:tagName`
- 6 Pull image:** `docker pull yourUserName/repoName:tagName`
- 7 Run a container on pulled image:** `docker run --name teammate -e MYSQL_ROOT_PASSWORD=password -p 3306:3306 -d yourUserName/repoName:tagName`

This only initializes the databases. Any changes will have to be shared through dumps.